


# CPU<sub>s</sub>

COEN 421/6341: Embedded Systems Design

1

## Programming Input and Output

- Input and output devices usually have some analog or nonelectronic component
  - e.g., disk drive has a rotating disk and analog read/write electronics
- Typical digital interface to CPU



- **Status reg.:** provide information about the device's operation (e.g., whether the current transaction has completed)
- **Data reg.:** hold values that are treated as data by the device (e.g., temperature)

COEN 421/6341: Embedded Systems Design

2

2

## Input and Output Primitives

- Microprocessors can provide programming support for input and output in two ways
  - 1) I/O instructions
    - Some architectures, such as the Intel x86, provide special instructions (in and out in the case of the Intel x86) for input and output
    - **IN** and **OUT** transfer data between an I/O device and the microprocessor's accumulator (AL, AX or EAX)
    - **IN AL, 19H**; 8-bits are saved to AL from I/O port 19H

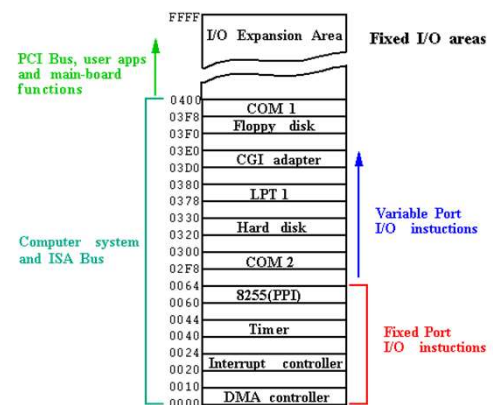
COEN 421/6341: Embedded Systems Design

3

3

## Input and Output Primitives

- Microprocessors can provide programming support for input and output in two ways
  - 2) Memory-mapped I/O
    - Provides addresses for the registers in each I/O device
    - It is the most common way to implement I/O is by memory mapping



COEN 421/6341: Embedded Systems Design

4

4

## I/O Primitives / MMIO

- ARM Cortex-M processors use MMIO to access peripheral registers
- Peripheral registers are mapped to a small memory region:
  - Starting at 0x40000000 on STM32L4

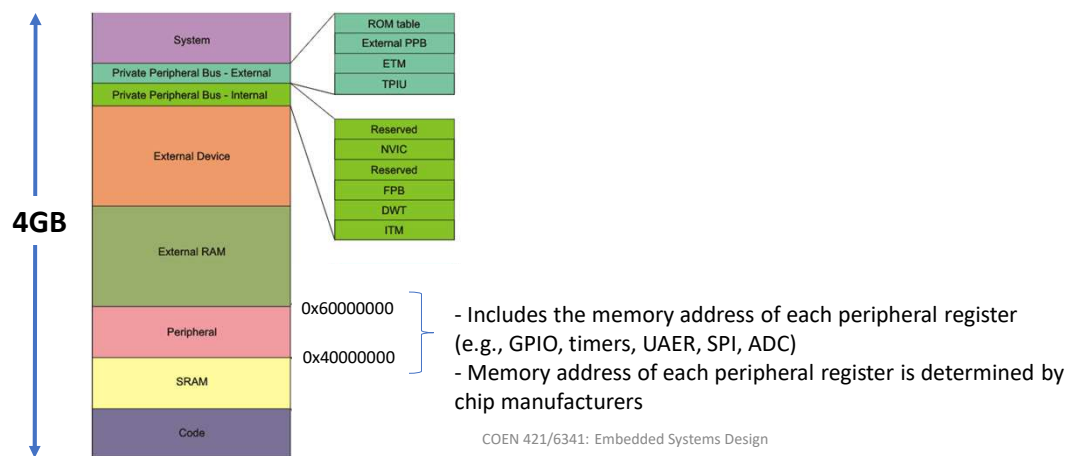
COEN 421/6341: Embedded Systems Design

5

5

## I/O Primitives / MMIO

- ARM Cortex-M processors:



COEN 421/6341: Embedded Systems Design

6

6

## I/O Primitives / MMIO

- ARM Cortex-M processors use MMIO to access peripheral registers
- Peripheral registers are mapped to a small memory region:
  - Starting at 0x40000000 on STM32L4
    - Includes the memory addresses of all on-chip peripherals:
      - GPIO
      - Timers
      - UART
      - SPI
      - ADC
      - ...

COEN 421/6341: Embedded Systems Design

7

7

## I/O Primitives / MMIO

- ARM Cortex-M processors use MMIO to access peripheral registers
- Peripheral registers are mapped to a small memory region:
  - Starting at 0x40000000 on STM32L4
    - Includes the memory addresses of all on-chip peripherals:
- Address of each peripheral register:
  - Determined by chip manufacturers

COEN 421/6341: Embedded Systems Design

8

8

## I/O Primitives / MMIO

- Programming/Interacting with I/O peripherals
  - Find out the address where the peripheral is memory mapped
    - Check the microcontroller datasheet, which provides a memory map
    - e.g., STM32L4

0x5FFF FFFF	Reserved
0x5006 0C00	AHB2
0x4800 0000	Reserved
0x4002 4400	AHB1
0x4002 0000	Reserved
0x4001 6400	APB2
0x4001 0000	Reserved
0x4000 9800	APB1
0x4000 0000	

COEN 421/6341: Embedded Systems Design

9

9

## I/O Primitives / MMIO

- Programming/Interacting with I/O peripherals
  - Find out the address where the peripheral is memory mapped
    - Check the microcontroller datasheet, which provides a memory map
    - e.g., STM32L4

0x5FFF FFFF	Reserved
0x5006 0C00	AHB2
0x4800 0000	Reserved
0x4002 4400	AHB1
0x4002 0000	Reserved
0x4001 6400	APB2
0x4001 0000	Reserved
0x4000 9800	APB1
0x4000 0000	

\*just part of it

Bus	Boundary address	Size (bytes)	Peripheral	Peripheral register map
APB1	0x4000 5000 - 0x4000 53FF	1 KB	UART5	Section 40.8.12: USART register map
	0x4000 4C00 - 0x4000 4FFF	1 KB	UART4	Section 40.8.12: USART register map
	0x4000 4800 - 0x4000 4BFF	1 KB	USART3	Section 40.8.12: USART register map
	0x4000 4400 - 0x4000 47FF	1 KB	USART2	Section 40.8.12: USART register map
	0x4000 4000 - 0x4000 43FF	1 KB	Reserved	-
	0x4000 3C00 - 0x4000 3FFF	1 KB	SPI3	Section 42.6.8: SPI register map
	0x4000 3800 - 0x4000 3BFF	1 KB	SPI2	Section 42.6.8: SPI register map
	0x4000 3400 - 0x4000 37FF	1 KB	Reserved	-
	0x4000 3000 - 0x4000 33FF	1 KB	IWDG	Section 36.4.6: IWDG register map
	0x4000 2C00 - 0x4000 2FFF	1 KB	WWDG	Section 37.5.4: WWDG register map
	0x4000 2800 - 0x4000 2BFF	1 KB	RTC	Section 38.6.21: RTC register map
	0x4000 2400 - 0x4000 27FF	1 KB	LCD <sup>(4)</sup>	Section 25.6.6: LCD register map
	0x4000 1800 - 0x4000 2400	3 KB	Reserved	-
	0x4000 1400 - 0x4000 17FF	1 KB	TIM7	Section 33.4.9: TIMx register map
	0x4000 1000 - 0x4000 13FF	1 KB	TIM6	Section 33.4.9: TIMx register map
	0x4000 0C00 - 0x4000 0FFF	1 KB	TIM5	Section 31.4.26: TIMx register map
	0x4000 0800 - 0x4000 0BFF	1 KB	TIM4	Section 31.4.26: TIMx register map
	0x4000 0400 - 0x4000 07FF	1 KB	TIM3	Section 31.4.26: TIMx register map
	0x4000 0000 - 0x4000 03FF	1 KB	TIM2	Section 31.4.26: TIMx register map

COEN 421/6341: Embedded Systems Design

10

10

## I/O Primitives / MMIO

- Programming/Interacting with I/O peripherals
  - Find out the address where the peripheral is memory mapped
    - Check the microcontroller datasheet, which provides a memory map
  - Find out the exact register address in order to program the registers
    - Check the peripheral datasheet
      - Control registers
      - Data registers
      - status registers
      - ...

COEN 421/6341: Embedded Systems Design

11

11

## I/O Primitives / MMIO

- Programming/Interacting with I/O peripherals
  - Find out the address where the peripheral is memory mapped
    - Check the microcontroller datasheet, which provides a memory map
  - Find out the exact register address in order to program the registers
    - e.g., STM32L4 GPIOA registers

Bus	Boundary address	Size (bytes)	Peripheral	Peripheral register map
AHB4	0xA000 1000 - 0xA000 13FF	1 KB	QUADSPI	<a href="#">Section 17.6.14: QUADSPI register map</a>
AHB3	0xA000 0400 - 0xA000 0FFF	3 KB	Reserved	-
AHB3	0xA000 0000 - 0xA000 03FF	1 KB	FMC	<a href="#">Section 16.7.8: FMC register map</a>
-	0x5006 0C00 - 0x5FFF FFFF	~260 MB	Reserved	-
-	0x5006 0800 - 0x5006 0BFF	1 KB	RNG	<a href="#">Section 27.7.4: RNG register map</a>
-	0x5006 0400 - 0x5006 07FF	1 KB	HASH	<a href="#">Section 29.7.8: HASH register map</a>
-	0x5006 0000 - 0x5006 03FF	1 KB	AES <sup>(2)</sup>	<a href="#">Section 28.7.18: AES register map</a>
-	0x5005 0400 - 0x5005 FFFF	63 KB	Reserved	-
-	0x5005 0000 - 0x5005 03FF	1 KB	DCMI	<a href="#">Section 20.5.12: DCMI register map</a>
-	0x5004 0400 - 0x5004 FFFF	63 KB	Reserved	-
-	0x5004 0000 - 0x5004 03FF	1 KB	ADC	<a href="#">Section 18.8: ADC register map</a>
-	0x5000 0000 - 0x5003 FFFF	256 KB	OTG_FS	<a href="#">Section 47.15.57: OTG_FS register map</a>
AHB2	0x4800 2400 - 0x4FFF FFFF	~127 MB	Reserved	-
AHB2	0x4800 2000 - 0x4800 23FF	1 KB	GPIOI	<a href="#">Section 8.5.13: GPIO register map</a>
AHB2	0x4800 1C00 - 0x4800 1FFF	1 KB	GPIOH	<a href="#">Section 8.5.13: GPIO register map</a>
AHB2	0x4800 1800 - 0x4800 1BFF	1 KB	GPIOG	<a href="#">Section 8.5.13: GPIO register map</a>
AHB2	0x4800 1400 - 0x4800 17FF	1 KB	GPIOF	<a href="#">Section 8.5.13: GPIO register map</a>
AHB2	0x4800 1000 - 0x4800 13FF	1 KB	GPIOE	<a href="#">Section 8.5.13: GPIO register map</a>
AHB2	0x4800 0C00 - 0x4800 0FFF	1 KB	GPIOD	<a href="#">Section 8.5.13: GPIO register map</a>
AHB2	0x4800 0800 - 0x4800 0BFF	1 KB	GPIOC	<a href="#">Section 8.5.13: GPIO register map</a>
AHB2	0x4800 0400 - 0x4800 07FF	1 KB	GPIOB	<a href="#">Section 8.5.13: GPIO register map</a>
AHB2	0x4800 0000 - 0x4800 03FF	1 KB	GPIOA	<a href="#">Section 8.5.13: GPIO register map</a>
-	0x4002 BC00 - 0x47FF FFFF	~127 MB	Reserved	-

COEN 421/6341: Embedded Systems Design

12

12

## I/O Primitives / MMIO

Registers					Offset	Register name	31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0		
AHB4	0xA000 1000 - 0xA000 13FF	1 KB	QUADSPI	Section 17.6.14: QUADSPI map	0x00	GPIOA_MODER	MODE15[1:0]	MODE14[1:0]	MODE13[1:0]	MODE12[1:0]	MODE11[1:0]	MODE10[1:0]	MODE9[1:0]	MODE8[1:0]	MODE7[1:0]	MODE6[1:0]	MODE5[1:0]	MODE4[1:0]	MODE3[1:0]	MODE2[1:0]	MODE1[1:0]	MODE0[1:0]																		
							1	0	1	0	1	0	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
AHB3	0xA000 0400 - 0xA000 0FFF	3 KB	Reserved	-	0x04	GPIOx_OTYPER (where x = A..I)	OT15	OT14	OT13	OT12	OT11	OT10	OT9	OT8	OT7	OT6	OT5	OT4	OT3	OT2	OT1	OT0																		
-	0x5006 0C00 - 0x5FFF FFFF	~260 MB	Reserved	-			Reset value																	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
AHB2	0x5006 0800 - 0x5006 0BFF	1 KB	RNG	Section 27.7.4: RNG register	0x08	GPIOA_OSPEEDR	OSPEED15[1:0]	OSPEED14[1:0]	OSPEED13[1:0]	OSPEED12[1:0]	OSPEED11[1:0]	OSPEED10[1:0]	OSPEED9[1:0]	OSPEED8[1:0]	OSPEED7[1:0]	OSPEED6[1:0]	OSPEED5[1:0]	OSPEED4[1:0]	OSPEED3[1:0]	OSPEED2[1:0]	OSPEED1[1:0]	OSPEED0[1:0]																		
	0x5006 0400 - 0x5006 07FF	1 KB	HASH	Section 29.7.8: HASH register			Reset value	0	0	0	0	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0x5006 0000 - 0x5006 03FF	1 KB	AES <sup>(2)</sup>	Section 28.7.18: AES register	0x0C	GPIOA_PUPDR	PUPDR15[1:0]	PUPDR14[1:0]	PUPDR13[1:0]	PUPDR12[1:0]	PUPDR11[1:0]	PUPDR10[1:0]	PUPDR9[1:0]	PUPDR8[1:0]	PUPDR7[1:0]	PUPDR6[1:0]	PUPDR5[1:0]	PUPDR4[1:0]	PUPDR3[1:0]	PUPDR2[1:0]	PUPDR1[1:0]	PUPDR0[1:0]																		
	0x5005 0400 - 0x5005 FFFF	63 KB	Reserved	-			Reset value	0	1	1	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0x5005 0000 - 0x5005 03FF	1 KB	DCMI	Section 20.5.12: DCMI register	...																																			
	0x5004 0400 - 0x5004 FFFF	63 KB	Reserved	-																																				
	0x5004 0000 - 0x5004 03FF	1 KB	ADC	Section 18.8: ADC register																																				
	0x5000 0000 - 0x5003 FFFF	256 KB	OTG_FS	Section 47.15.57: OTG_FS map																																				
	0x4800 2400 - 0x4FFF FFFF	~127 MB	Reserved	-																																				
	0x4800 2000 - 0x4800 23FF	1 KB	GPIOI	Section 8.5.13: GPIO register map																																				
	0x4800 1C00 - 0x4800 1FFF	1 KB	GPIOH	Section 8.5.13: GPIO register map																																				
	0x4800 1800 - 0x4800 1BFF	1 KB	GPIOG	Section 8.5.13: GPIO register map																																				
	0x4800 1400 - 0x4800 17FF	1 KB	GPIOF	Section 8.5.13: GPIO register map																																				
	0x4800 1000 - 0x4800 13FF	1 KB	GPIOE	Section 8.5.13: GPIO register map																																				
	0x4800 0C00 - 0x4800 0FFF	1 KB	GPIOD	Section 8.5.13: GPIO register map																																				
	0x4800 0800 - 0x4800 0BFF	1 KB	GPIOC	Section 8.5.13: GPIO register map																																				
0x4800 0400 - 0x4800 07FF	1 KB	GPIOB	Section 8.5.13: GPIO register map																																					
0x4800 0000 - 0x4800 03FF	1 KB	GPIOA	Section 8.5.13: GPIO register map																																					
-	0x4002 BC00 - 0x47FF FFFF	~127 MB	Reserved	-																																				

CCN-41/6341: Embedded Systems Design

13

COEN 421/6341: Embedded Systems Design

13

13

## I/O Primitives / MMIO

### Example 3.1 Memory-Mapped I/O on ARM

We can use the EQU pseudo-op to define a symbolic name for the memory location of our I/O device:

```
DEV1    EQU 0x1000
```

Given that name, we can use the following standard code to read and write the device register:

```
LDR r1, #DEV1      ; set up device address
LDR r0, [r1]        ; read DEV1
LDR r0, #8          ; set up value to write

STR r0, [r1]        ; write 8 to device
```

COEN 421/6341: Embedded Systems Design

14

14

## I/O Primitives / MMIO

- How can we directly write I/O devices in a high-level language such as C?
  - We can use pointers to manipulate addresses of I/O devices
  - **peek** and **poke** functions that read and write arbitrary memory locations

```
#define DEV1 0x1000
...
dev_status = peek(DEV1); /* read device register */
...
poke(DEV1,8); /* write 8 to device register */
```

COEN 421/6341: Embedded Systems Design

15

15

## I/O Primitives / MMIO

- How can we directly write I/O devices in a high-level language such as C?

```
int peek(char *location) {
    return *location; /* de-reference location pointer */
}

void poke(char *location, char newval) {
    (*location) = newval; /* write to location */
}
```

COEN 421/6341: Embedded Systems Design

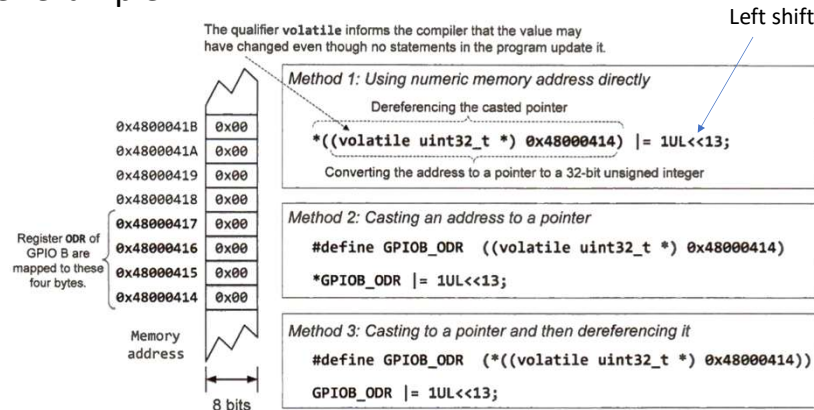
16

16



## I/O Primitives / MMIO

- Another example:



COEN 421/6341: Embedded Systems Design

17

17

## Busy-Wait I/O

- Simplest way to communicate with devices
- Devices are typically slower than the CPU and may require many cycles to complete an operation
- CPU must wait for one I/O operation to complete before starting the next one when performing multiple operations on a single device
- **Polling**
  - It is the process of asking an I/O device whether it is finished
  - This is done by reading the device's status register

COEN 421/6341: Embedded Systems Design

18

18

## Busy-Wait I/O

```
#define OUT_CHAR 0x1000 /* output device character register */
#define OUT_STATUS 0x1001 /* output device status register */

char *mystring = "Hello, world." /* string to write */
char *current_char; /* pointer to current position in string */
current_char = mystring; /* point to head of string */
while (*current_char != '\0') { /* until null character */
    poke(OUT_CHAR, *current_char); /* send character to device */
    while (peek(OUT_STATUS) != 0); /* keep checking status */
    current_char++; /* update character pointer */
}
```

COEN 421/6341: Embedded Systems Design

19

19

## Busy-Wait I/O

- Copying Characters from Input to Output using Busy-Wait I/O

```
#define IN_DATA 0x1000
#define IN_STATUS 0x1001
#define OUT_DATA 0x1100
#define OUT_STATUS 0x1101
```

- Input device status register
  - 1: new character is ready
  - 0: new character has been read
- Output device status register
  - 1: start writing the character
  - 0: writing operation has been concluded

COEN 421/6341: Embedded Systems Design

20

20

## Busy-Wait I/O

- Copying Characters from Input to Output using Busy-Wait I/O

```
while (TRUE) { /* perform operation forever */
    /* read a character into achar */
    while (peek(IN_STATUS) == 0); /* wait until ready */
    achar = (char)peek(IN_DATA); /* read the character */
    /* write achar */
    poke(OUT_DATA, achar);
    poke(OUT_STATUS, 1); /* turn on device */
    while (peek(OUT_STATUS) != 0); /* wait until done */
}
```

COEN 421/6341: Embedded Systems Design

21

21

## Interrupts


- Busy-wait I/O is extremely inefficient
  - CPU does nothing but test the device status while the I/O transaction is in progress
- Interrupt
  - Allows the processor to perform multiple tasks “simultaneously”
  - Allows devices to signal the CPU and to force execution of a particular piece of code
  - When an interrupt occurs:
    - The interrupt mechanism saves the value of the PC
    - The program counter’s value is changed to point to an **interrupt handler** routine (device driver)

COEN 421/6341: Embedded Systems Design

22

22

## Interrupts



- The I/O device's logic decides when to interrupt
- The CPU may not be able to immediately service an interrupt request because it may be doing something else that must be finished first
- When CPU acknowledge the interrupt, it changes the PC to point to the device's handler
- The interrupt handler operates much like a subroutine, except that it is not called by the executing program

COEN 421/6341: Embedded Systems Design

23

23

## Interrupts

- The CPU checks the interrupt request line at the beginning of execution of every instruction
- The CPU handles the interrupt if the interrupt request has been asserted
- The CPU does not fetch the instruction pointed to by the PC
  - CPU sets the PC to a predefined location (i.e., the beginning of the interrupt handling – interrupt service routine)

COEN 421/6341: Embedded Systems Design

24

24

## Interrupt Service Routine

- Special subroutine
- Invoked automatically by the hardware in response to an interrupt
- Each ISR has a default implementation in the system startup code
- There is an ISR associated with each type of interrupt
- ISRs are stored in a special memory part
- Interrupt vector table

COEN 421/6341: Embedded Systems Design

25

25

## Copying Characters from I to O with Basic Interrupts

- Assumption:
  - Write C functions that act as interrupt handlers
  - Functions will read and write status and data registers
  - Use of a global variable ***achar*** for the input handler to pass the character to the foreground program
  - Because the foreground program does not know when an interrupt occurs, we also use a global Boolean variable, ***gotchar***, to signal when a new character has been received
    - Foreground program: the program that runs when no interrupt is being handled

COEN 421/6341: Embedded Systems Design

26

26

## Copying Characters from I to O with Basic Interrupts

- Handlers

```
void input_handler() { /* get a character and put in global */
    achar = peek(IN_DATA); /* get character */
    gotchar = TRUE; /* signal to main program */
    poke(IN_STATUS, 0); /* reset status to initiate next
transfer */
}
void output_handler() { /* react to character being sent */
    /* don't have to do anything */
}
```

COEN 421/6341: Embedded Systems Design

27

27

## Copying Characters from I to O with Basic Interrupts

- Main program

```
main() {
    while (TRUE) { /* read then write forever */
        if (gotchar) { /* write a character */
            poke(OUT_DATA, achar); /* put character in device */
            poke(OUT_STATUS, 1); /* set status to initiate write */
            gotchar = FALSE; /* reset flag */
        }
    }
}
```

COEN 421/6341: Embedded Systems Design

28

28

## Copying Characters from I to O with Basic Interrupts

- The main program is somewhat simpler when interrupt handler was used
- However, it still does not let the foreground program do useful work
- Goal: Perform reads and writes independently
- Elastic buffer to hold the characters
  - It will allow the input and output devices to run at different rates


COEN 421/6341: Embedded Systems Design

29

29

## Copying Characters from I to O with Basic Interrupts

- Characters are added to the **tail** when received
- Characters are taken from the **head** when ready for output
- The head and tail wrap around the end of the buffer to make most efficient use of it




COEN 421/6341: Embedded Systems Design

30

30

## Copying Characters from I to O with Basic Interrupts

- First character is ready and added




COEN 421/6341: Embedded Systems Design

31

31

## Copying Characters from I to O with Basic Interrupts

- One character in the buffer is left unused to make possible the distinguishing buffer full of buffer empty



COEN 421/6341: Embedded Systems Design


32

32



## Copying Characters from I to O with Basic Interrupts

- Output goes past the end of io\_buf



COEN 421/6341: Embedded Systems Design

33

33

## Copying Characters from I to O with Basic Interrupts

- Assume two interrupt handling:
  - 1 for the input device
  - 1 for the output device

```
void full_buffer() { /* returns TRUE if buffer is full */
    (buf_tail+1) % BUF_SIZE == buf_head;
}
```

```
#define IN_DATA 0x1000
#define IN_STATUS 0x1001
void input_handler() {
    char achar;
    if (full_buffer()) /* error */
        error = 1;
    else { /* read the character and update pointer */
        achar = peek(IN_DATA); /* read character */
        add_char(achar); /* add to queue */
    }
    poke(IN_STATUS, 0); /* set status register back to 0 */
    /* if buffer was empty, start a new output transaction */
    if (nchars() == 1) { /* buffer had been empty until this
        interrupt */
        poke(OUT_DATA, remove_char()); /* send character */
        poke(OUT_STATUS, 1); /* turn device on */
    }
}
```

COEN 421/6341: Embedded Systems Design


34

34

## Copying Characters from I to O with Basic Interrupts

- Full buffer (e.g.)

```
void full_buffer() { /* returns TRUE if buffer is full */
    (buf_tail+1) % BUF_SIZE == buf_head ;
}
```



COEN 421/6341: Embedded Systems Design

35

35

## Copying Characters from I to O with Basic Interrupts

- Assume two interrupt handling:

- 1 for the input device
- 1 for the output device

```
void add_char(char achar) { /* add a character to the buffer
head */
    io_buf[buf_tail++] = achar;
    /* check pointer */
    if (buf_tail == BUF_SIZE)
        buf_tail = 0;
}
```

```
#define IN_DATA 0x1000
#define IN_STATUS 0x1001
void input_handler() {
    char achar;
    if (full_buffer()) /* error */
        error = 1;
    else { /* read the character and update pointer */
        achar = peek(IN_DATA); /* read character */
        add_char(achar); /* add to queue */
    }
    poke(IN_STATUS, 0); /* set status register back to 0 */
    /* if buffer was empty, start a new output transaction */
    if (nchars() == 1) { /* buffer had been empty until this
interrupt */
        poke(OUT_DATA, remove_char()); /* send character */
        poke(OUT_STATUS, 1); /* turn device on */
    }
}
```

COEN 421/6341: Embedded Systems Design

36

36

## Copying Characters from I to O with Basic Interrupts

- Assume two interrupt handling:

- 1 for the input device
- 1 for the output device

It means a new character is being read and added in the buffer

```
#define IN_DATA 0x1000
#define IN_STATUS 0x1001
void input_handler() {
    char achar;
    if (full_buffer()) /* error */
        error = 1;
    else { /* read the character and update pointer */
        achar = peek(IN_DATA); /* read character */
        add_char(achar); /* add to queue */
    }
    poke(IN_STATUS, 0); /* set status register back to 0 */
    /* if buffer was empty, start a new output transaction */
    if (nchars() == 1) { /* buffer had been empty until this
        interrupt */
        poke(OUT_DATA, remove_char()); /* send character */
        poke(OUT_STATUS, 1); /* turn device on */
    }
}
```

COEN 421/6341: Embedded Systems Design

37

37

## Copying Characters from I to O with Basic Interrupts

- Assume two interrupt handling:

- 1 for the input device
- 1 for the output device

```
int nchars() { /* returns the number of characters in the
buffer */
    if (buf_head >= buf_tail) return buf_head - buf_tail;
    else return BUF_SIZE - buf_tail - buf_head;
}
```

Just when the first character is added in the buffer

```
#define IN_DATA 0x1000
#define IN_STATUS 0x1001
void input_handler() {
    char achar;
    if (full_buffer()) /* error */
        error = 1;
    else { /* read the character and update pointer */
        achar = peek(IN_DATA); /* read character */
        add_char(achar); /* add to queue */
    }
    poke(IN_STATUS, 0); /* set status register back to 0 */
    /* if buffer was empty, start a new output transaction */
    if (nchars() == 1) { /* buffer had been empty until this
        interrupt */
        poke(OUT_DATA, remove_char()); /* send character */
        poke(OUT_STATUS, 1); /* turn device on */
    }
}
```

COEN 421/6341: Embedded Systems Design

38

38

## Copying Characters from I to O with Basic Interrupts

- Assume two interrupt handling: 1 for the input device + 1 for the output device

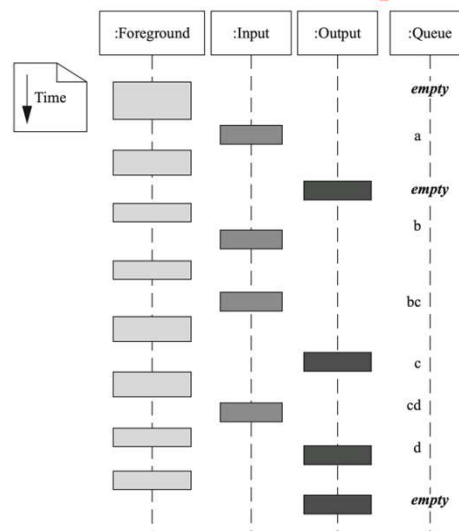
```
#define OUT_DATA 0x1100
#define OUT_STATUS 0x1101
void output_handler() {
    if (!empty_buffer()) { /* start a new character */
        poke(OUT_DATA, remove_char()); /* send character */
        poke(OUT_STATUS, 1); /* turn device on */
    }
}
```

COEN 421/6341: Embedded Systems Design

39

39

## Copying Characters from I to O with Basic Interrupts



40

40

## Interrupts: Priorities and Vectors

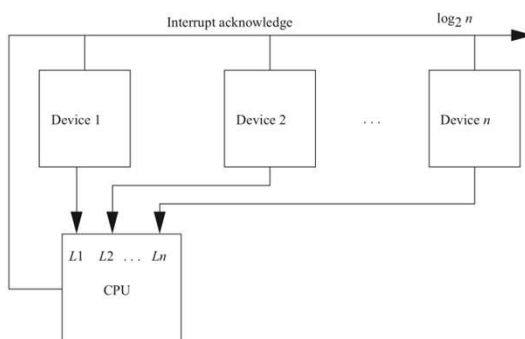
- Most embedded systems have more than one I/O device
- There is a need for a mechanism for allowing multiple devices to interrupt
- Two approaches can be used to handle multiple I/O devices
  - **Interrupt Priorities:**
    - Allow the CPU to recognize some interrupts as more important than others
  - **Interrupt vectors:**
    - Allow the interrupting device to specify its handle

COEN 421/6341: Embedded Systems Design

41

41

## Interrupt Priorities



- CPU provides several different interrupt request signals ( $L1, L2, \dots, L_n$ )
- Typically, the lower-numbered interrupt lines are given higher priority
- Most CPUs use a set of signals that provide the priority number of the winning interrupt in binary form
- A device knows that its interrupt request was accepted by seeing its own priority number on the interrupt acknowledge lines

COEN 421/6341: Embedded Systems Design

42

42

## Interrupt Priorities

- Masking:
  - Mechanism used to ensure that a lower-priority interrupt does not occur when a higher-priority interrupt is being handled
  - CPU stores in an internal register the priority level of that interrupt, when it acknowledge an interrupt
  - Priority of subsequent interrupts are checked against the priority register
    - New request is acknowledged only if it has higher priority than the currently pending interrupt
  - Priority register is reset when the interrupt handler exits

COEN 421/6341: Embedded Systems Design

43

43

## Interrupt Priorities

- Nonmaskable interrupt (NMI)
  - The highest-priority interrupt
  - The NMI cannot be turned off
  - Usually reserved for interrupts caused by power failures
    - A simple circuit can be used to detect a dangerously low power supply, and the NMI interrupt handler can be used to save critical state in nonvolatile memory, turn off I/O devices to eliminate spurious device operation during power- down, and so on


COEN 421/6341: Embedded Systems Design

44

44

## Interrupt Priorities

- Most CPUs provide a relatively small number of interrupt priority levels, such as eight
- When several devices naturally assume the same priority:



- The CPU will call the interrupt handler associated with this priority
- That handler does not know which of the devices actually requested the interrupt
- The handler uses software polling to check the status of each device:
  - e.g., it would read the status registers of 1, 2, and 3 to see which of them is ready and requesting service

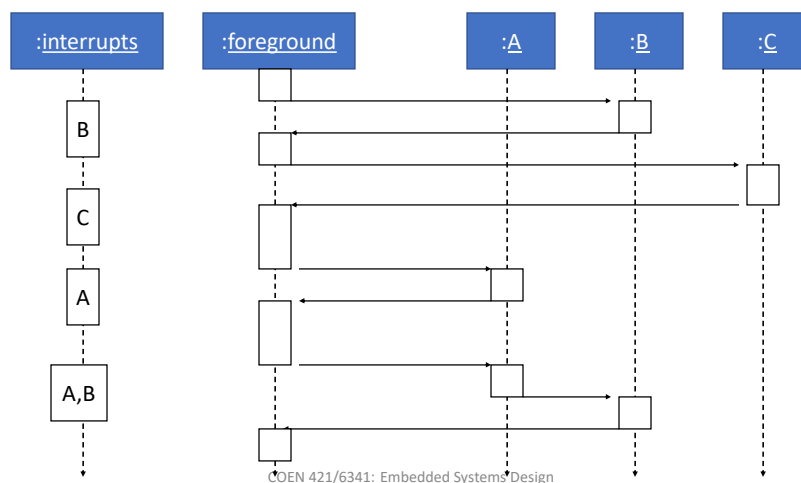
COEN 421/6341: Embedded Systems Design

45

45

## Interrupt Priorities

A has priority 1 (highest priority), B has priority 2, and C has priority 3




COEN 421/6341: Embedded Systems Design

46

46

## Interrupt Vectors

- Provides the ability to define the interrupt handler that should service a request from a device



Interrupt vector table head

Handler 1
Handler 3
Handler 4
Handler 2

Interrupt vector table

Vector 0

Vector 1

Vector 2

Vector 3

- Interrupt vector lines run from the devices to the CPU
- After a device's request is acknowledged, it sends its interrupt vector over those lines to the CPU
- The CPU then uses the vector number as an index in a table stored in memory
- The location referenced in the interrupt vector table by the vector number gives the address of the handler

COEN 421/6341: Embedded Systems Design

47

47

## Interrupts

- Most modern CPUs implement both prioritized and vectored interrupts
- Priorities determine which device is serviced first, and vectors determine what routine is used to service the interrupt

COEN 421/6341: Embedded Systems Design

48

48



## The Interrupt Handling Process

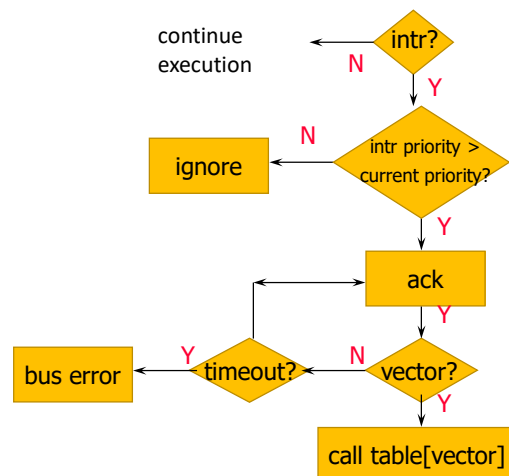
- Device requests an interrupt
- The CPU checks for pending interrupts at the beginning of an instruction
  - It will acknowledge the highest-priority pending interrupt (which has a higher priority than that given in the interrupt priority register)
- The device receives the acknowledgment and sends the CPU its interrupt vector
- The CPU looks up the device handler address in the interrupt vector table using the vector as an index
- The device driver (interrupt handler routine) may save additional CPU state
- It then performs the required operations on the device
- It then restores any saved state and executes the interrupt return instruction
- The interrupt return instruction restores the PC and other automatically saved states to return execution to the code that was interrupted

COEN 421/6341: Embedded Systems Design

49

49

## The Interrupt Handling Process

COEN 421/6341: Embedded Systems Design  
Computers as Components 4e © 2016 Marilyn Wolf

50

50

## Cortex-M Interrupts

- 256 types of interrupts
- Each interrupt is identified by a unique number [-15, 240] (except reset interrupt)
- Interrupt numbers are defined by ARM and chip manufacturers collectively
- Interrupt numbers are divided into two groups
  - System interrupts (or system exceptions) [-15, -1]
    - Those exceptions (interrupts) come from the processor core
    - Defined by ARM
  - Peripheral interrupts [0, 240]
    - Defined by chip manufacturers
    - Total number of peripheral interrupts supported varies among chips

COEN 421/6341: Embedded Systems Design

51

51

## Cortex-M Interrupts

Cortex-M4 Processor  
Exceptions Numbers

-14	Non-maskable interrupt
-13	Hard fault
-12	Memory management
-11	Bus fault
-10	Usage fault
-5	Supervisor call (SVCall)
-4	Debug monitor
-2	PendSV
-1	SysTick

STM32L4 specific  
Interrupt Numbers

0	WWDG	16	DMA1_CH6	32	I2C1_ER	48	FMC	64	COMP	80	RNG
1	PVD	17	DMA1_CH7	33	I2C2_EV	49	SDMMC1	65	LPTIM1	81	FPU
2	TAMPER_STAMP	18	ADC1_ADC2	34	I2C2_ER	50	TIM5	66	LPTIM2		
3	RTC_WKUP	19	CAN1_TX	35	SPI1	51	SPI3	67	OTG_FS		
4	FLASH	20	CAN1_RX0	36	SPI2	52	UART4	68	DMA2_Channel6		
5	RCC	21	CAN1_RX1	37	USART1	53	UART5	69	DMA2_Channel7		
6	EXTI0	22	CAN1_SCE	38	USART2	54	TIM8_DAC	70	LP_UART1		
7	EXTI1	23	EXTI9_5	39	USART3	55	TIM7	71	QUADSPI		
8	EXTI2	24	TIM1_BRK	40	EXTI15_10	56	DMA2_Channel1	72	I2C3_EV		
9	EXTI3	25	TIM1_UP	41	RTC_Alarm	57	DMA2_Channel2	73	I2C3_ER		
10	EXTI4	26	TIM1_TRG	42	DFSDM3	58	DMA2_Channel3	74	SAI1		
11	DMA1_CH1	27	TIM1_CC	43	TIM8_BRK	59	DMA2_Channel4	75	SAI2		
12	DMA1_CH2	28	TIM2	44	TIM8_UP	60	DMA2_Channel5	76	SWPMI1		
13	DMA1_CH3	29	TIM3	45	TIM8_TRG	61	DFSDM0	77	TSC		
14	DMA1_CH4	30	TIM4	46	TIM8_CC	62	DFSDM1	78	LCD		
15	DMA1_CH5	31	I2C1_EV	47	ADC3	63	DFSDM2	79			


COEN 421/6341: Embedded Systems Design

52

52

## Cortex-M Interrupts

- The numbering scheme allows software to distinguish system exceptions and peripheral interrupts
- Interrupt number is stored in the PSR when an interrupt is processed



- ARM Cortex-M does not store interrupt numbers in two's complement
- How can we represent the negative numbers?
  - Interrupt number in PSR adds a positive offset of 15

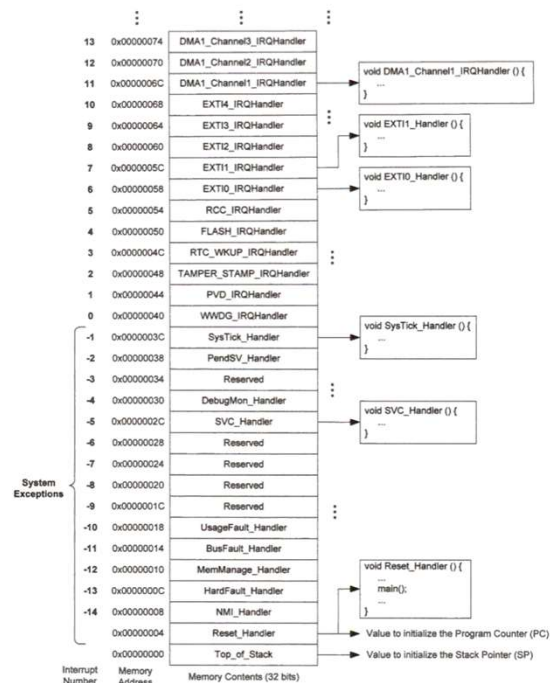
COEN 421/6341: Embedded Systems Design

53

53

## Cortex-M Interrupt Vector Table

- Interrupt Vector Table:
  - Special array used by Cortex-M to store the starting memory address of every ISR
- Address of ISR =  $IVT[i+15]$
- How can we find the memory address of an interrupt handler?

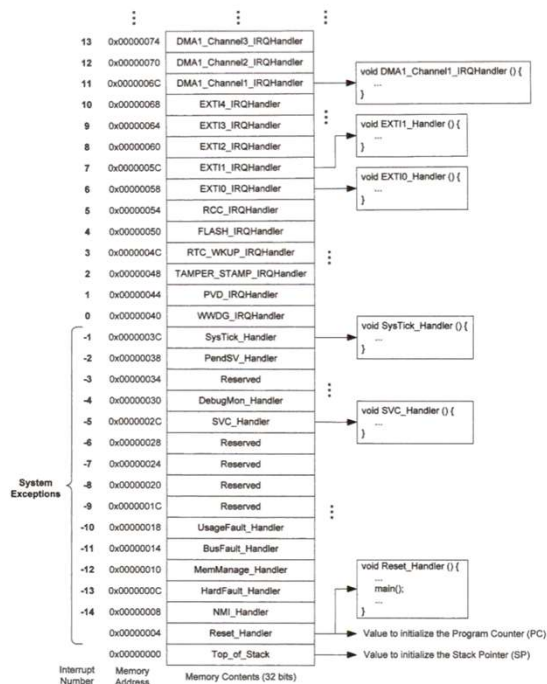


COEN 421/6341: Emb

54

## Cortex-M Interrupt Vector Table

- How can we find the memory address of an interrupt handler?
  - Example: SysTick\_Handler?
    - Interrupt vector table is stored at the memory address:
      - 0x00000004
    - SysTick interrupt is -1
      - $A = 0x00000004 + 4 \times (-1 + 15)$
      - $A = 0x00000003C$



COEN 421/6341: Emb

55

## Cortex-M Interrupts

- Enable specific interrupts and set their priority levels
- All interrupts are served based on their priority levels
- Priority levels -> given by interrupt number
  - A higher value of the interrupt priority number represents a lower priority
- A current interrupt being served is stopped if a higher priority interrupt occurs
- Processor resumes the low-priority interrupt when the handler of the higher priority interrupt completes


COEN 421/6341: Embedded Systems Design

56

56

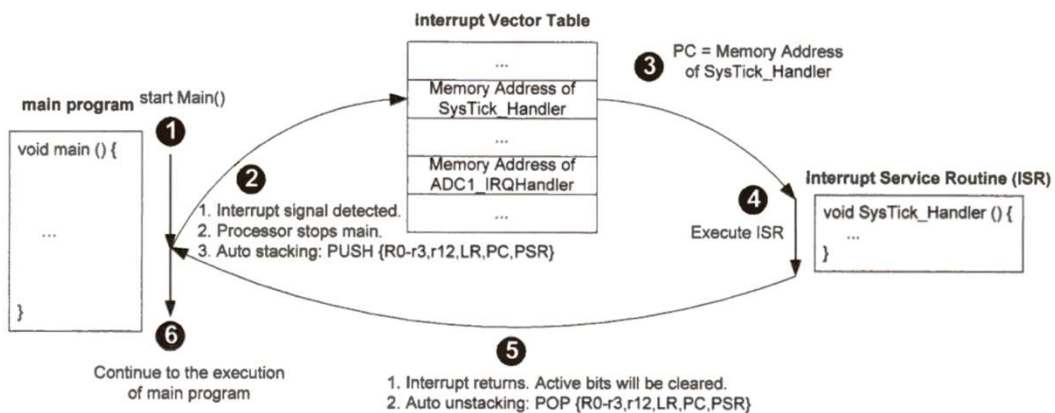
## Cortex-M Interrupts

- **Interrupt stacking:**
  - Before executing the interrupt handler, the current running environment must be preserved
  - 8 registers are preserved (r0, r1, r2, r3, r12, LR, PSR, and PC)
- **Interrupt Unstacking**
  - Recovers the environment that existed at the time instant immediately before the interrupt handler started to run
  - Pops the values of the r0, r1, r2, r3, r12, LR, PSR, and PC off the stack



57

## Cortex-M Interrupts



58

## Cortex-M Interrupts

- The interrupt service routine exits by running “BX LR”
- LR in an interrupt service routine indicates whether the processor uses the main stack or the process stack in the push and pop operations
  - LR shall be fixed to a special value to indicate whether the processor should unstack data out of the main stack (MSP) or process stack (PSP)

COEN 421/6341: Embedded Systems Design

59

59

## Nested Vectored Interrupt Controller

- Three key functions:
  - Enable and disable interrupts
  - Configure the preemption priority and sub-priority of a specific interrupt
  - Set and clear the handling bit of a specific interrupt
- Each interrupt has six control bits

Interrupt control bit	Corresponding register (32 bits)
Enable bit	Interrupt set enable register (ISER)
Disable bit	Interrupt clear enable register (ICER)
Pending bit	Interrupt set pending register (ISPR)
Un-pending bit	Interrupt clear pending register (ICPR)
Active bit	Interrupt active bit register (IABR)
Software trigger bit	Software trigger interrupt register (STIR)

COEN 421/6341: Embedded Systems Design

60

60

## Nested Vectored Interrupt Controller

- Three key functions:
  - Enable and disable interrupts
  - Configure the preemption priority and sub-priority of a specific interrupt
  - Set and clear the handling bit of a specific interrupt
- Each interrupt has six control bits and 8 registers for each control bit

Interrupt control bit	Corresponding register (32 bits)
Enable bit	Interrupt set enable register (ISER)
Disable bit	Interrupt clear enable register (ICER)
Pending bit	Interrupt set pending register (ISPR)
Un-pending bit	Interrupt clear pending register (ICPR)
Active bit	Interrupt active bit register (IABR)
Software trigger bit	Software trigger interrupt register (STIR)

COEN 421/6341: Embedded Systems Design

61

61

## Nested Vectored Interrupt Controller

- Interrupt control bit
  - Enable interrupt
    - Write an enable bit to 1 in ISER register
  - Disable interrupt
    - Write a disable bit to 1 in ICER register
  - Pending interrupt
    - If the microcontroller cannot process an interrupt immediately the corresponding pending bit is set in ISPR
    - If the pending interrupt is handled, the corresponding interrupt is removed from the pending list by setting the un-pending corresponding bit (ICPR)
  - Active interrupt
    - When an interrupt is disabled but its pending bit has already been set, this interrupt instance remains active and it is serviced before it is disabled

Interrupt control bit	Corresponding register (32 bits)
Enable bit	Interrupt set enable register (ISER)
Disable bit	Interrupt clear enable register (ICER)
Pending bit	Interrupt set pending register (ISPR)
Un-pending bit	Interrupt clear pending register (ICPR)
Active bit	Interrupt active bit register (IABR)
Software trigger bit	Software trigger interrupt register (STIR)

COEN 421/6341: Embedded Systems Design

62

62





## Nested Vectored Interrupt Controller

- Setting the proper bit in register ISER to enable interrupt IRQn

```
WordOffset = IRQn >> 5;           // Word Offset = IRQn/32
BitOffset = IRQn & 0x1F;         // Bit Offset = IRQn mod 32
NVIC->ISER[WordOffset] = 1 << BitOffset; // Enable interrupt
```

- E.g., IRQn = 50 (Timer5 interrupt)

```
NVIC->ISER[1] = 1 << 18;           // Enable Timer 5 interrupt
```

- $IRQn = j + 32 \times i$ 
  - $i = \text{floor}(\text{interrupt number}/32) \Rightarrow \text{floor}(50/32) = 1$
  - $j = \text{interrupt number} \% 32 \Rightarrow 50 \% 32 = 18$

COEN 421/6341: Embedded Systems Design

65

65

## Nested Vectored Interrupt Controller

- Assembly implementation of enabling or disabling a peripheral interrupt

```
; Input arguments:
;   r0: interrupt number of a peripheral interrupt
;   r1: 1 = Enable, 0 = Disable

Peripheral_Interrupt_Enable PROC
    PUSH {r4, lr}
    AND  r2, r0, #0x1F      ; Bit offset in a word
    MOV  r3, #1
    LSL  r3, r3, r2          ; r3 = 1 << (IRQn & 0x1F)
    LDR  r4, =NVIC_BASE

    CMP  r1, #0             ; Check whether enable or disable
    LDRNE r1, =NVIC_ISER0   ; Enable register base address
    LDREQ r1, =NVIC_ICER0   ; Disable register base address

    ADD  r1, r4, r1          ; r1 = addr. of NVIC->ISER0 or NVIC->ICER0
    LSR  r2, r0, #5          ; Memory offset (in words): IRQn >> 5
    LSL  r2, r2, #2          ; Calculate byte offset
    STR  r3, [r1, r2]        ; Enable/Disable interrupt
    POP  {r4, pc}
ENDP
```

66

66

## Nested Vectored Interrupt Controller

- Setting the proper bit in register **ICER** to enable interrupt IRQn

```
WordOffset = IRQn >> 5;           // WordOffset = IRQn/32
BitOffset = IRQn & 0x1F;          // BitOffset = IRQn mod 32
NVIC->ICER[WordOffset] = 1 << BitOffset; // Disable interrupt
```


COEN 421/6341: Embedded Systems Design

67

67

## Memory System Mechanism

- Caches
  - Small, fast memory that holds copies of some of the contents of main memory
  - Used to speed up reads and writes in memory systems
- Cache Controller:
  - Mediates between the CPU and the memory system comprised of the cache and main memory
  - Sends a memory request to the cache and main memory
- Cache Hit:
  - If the requested location is in the cache, the cache controller forwards the location's contents to the CPU and aborts the main memory request
- Cache Miss:
  - If the location is not in the cache, the controller waits for the value from main memory and forwards it to the CPU



COEN 421/6341: Embedded Systems Design

68

68

## Cache

- We can classify cache misses into several types depending on the situation that generated them
  - Compulsory miss (also known as a cold miss)
    - Occurs the first time a location is used
  - Capacity miss
    - Caused by a too-large working set
  - Conflict miss
    - Happens when two locations map to the same location in the cache.

COEN 421/6341: Embedded Systems Design

69

69

## Memory System Performance

- $t_{\text{cache}}$  : access time of the cache (a few nanoseconds)
- $t_{\text{main}}$  : main memory access time (~50 to 75ns)
- $t_{\text{cache}} \ll t_{\text{main}}$
- $h$ : hit ratio
- $t_{\text{av}} = ht_{\text{cache}} + (1-h)t_{\text{main}}$


COEN 421/6341: Embedded Systems Design

70

70

## Cache

- Modern CPUs may use multiple levels of cache
- L1: First-level cache is closest to the CPU
- L2: Second-level cache feeds the first-level cache
- ... and so on



$$t_{av} = h_1 t_{L1} + (h_2 - h_1) t_{L2} + (1 - h_2) t_{main}$$


COEN 421/6341: Embedded Systems Design

71

71

## Direct-Mapped Cache

- The cache consists of blocks
- Each line includes:
  - A **tag** to show which memory location is represented by this block
  - A **data** field that holds the content of that memory location
  - A **valid tag** to show whether the contents of this cache block are valid



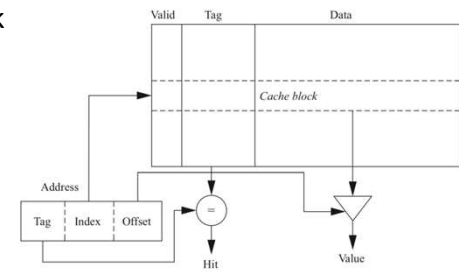
COEN 421/6341: Embedded Systems Design

72

72

## Direct-Mapped Cache

- Address: divided into 3 sections
  - **Index**: used to select which cache block to check
  - **Tag** is compared against the tag value in the block selected by the index
  - **Offset** is used to select the required value from the data field when the length of the data field is longer than the minimum addressable unit



COEN 421/6341: Embedded Systems Design

73

73

## Direct-Mapped Cache

- Write operation: We have to update the main memory as well
  - Write-through
    - Every write changes both the cache and the corresponding main memory location
  - Write-back
    - Write only when we remove a location from the cache


COEN 421/6341: Embedded Systems Design

74

74

## Direct-Mapped Cache

- Example: Assume a main memory of 16Kb and cache of 256bytes organized in 4byte blocks
- Memory address: 14-bits
- Cache lines:  $256/4 = 64$ :
  - Thus, 6 bits to address the lines
- 4 blocks in each cache line:
  - Thus, 2 bits to address the block
- 6 bits as tag



COEN 421/6341: Embedded Systems Design

75

75

## Set-Associative Cache

- Characterized by the number of banks or ways
- A set is formed by all the blocks (one for each bank) that share the same index
- Each set is implemented with a direct-mapped cache
- A cache request is broadcast to all banks (or ways) simultaneously
- If any of the ways has the location, the cache reports a hit
- The set-associative cache generally provides higher hit rates than the direct-mapped cache because conflicts between a small set of locations can be resolved within the cache

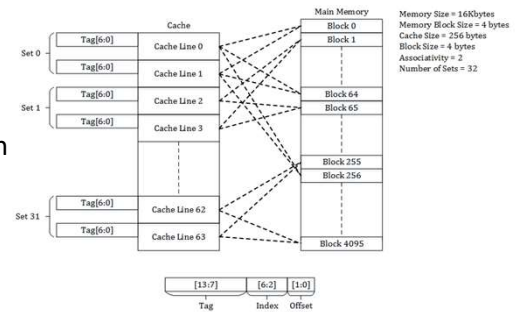
COEN 421/6341: Embedded Systems Design

76

76

## Set-Associative Cache

- Example: Assume a main memory of 16Kb and a 2-way set-associative of 256 bytes organized with a block size of 4
- Memory address: 14-bits long
- Cache:
  - 2-way set associative: means 2 cache lines in each set
  - Number of banks:  $256/(4*2) = 32$
- Addressing:
  - 2 bits to address the block
  - 5 bits to address the set
  - 7 bits for tag



COEN 421/6341: Embedded Systems Design

77

77

## Example

- Let's assume a memory where the address is given by 3 bits
- Data requests: 001, 010, 011, 100, 101, and 111

After 001 access:

Block	Tag	Data
00	—	—
01	0	1111
10	—	—
11	—	—

After 010 access:

Block	Tag	Data
00	—	—
01	0	1111
10	0	0000
11	—	—

After 011 access:

Block	Tag	Data
00	—	—
01	0	1111
10	0	0000
11	0	0110

Address	Data
000	0101
001	1111
010	0000
011	0110
100	1000
101	0001
110	1010
111	0100

COEN 421/6341: Embedded Systems Design

78

78

## Example

- Let's assume a memory where the address is given by 3 bits
- Data requests: 001, 010, 011, 100, 101, and 111

After 100 access (notice that the tag bit for this entry is 1):

Block	Tag	Data
00	1	1000
01	0	1111
10	0	0000
11	0	0110

After 101 access (overwrites the 01 block entry):

Block	Tag	Data
00	1	1000
01	1	0001
10	0	0000
11	0	0110

After 111 access (overwrites the 11 block entry):

Block	Tag	Data
00	1	1000
01	1	0001
10	0	0000
11	1	0100

Address	Data
000	0101
001	1111
010	0000
011	0110
100	1000
101	0001
110	1010
111	0100

COEN 421/6341: Embedded Systems Design

79

79

## Example: Set-Associative Cache

To start, let's consider that each bank has the same size of the example of the original direct-mapped cache

Data requests: 001, 010, 011, 100, 101, and 111

Block	Bank 0 tag	Bank 0 data	Bank 1 tag	Bank 1 data
00	1	1000	—	—
01	0	1111	1	0001
10	0	0000	—	—
11	0	0110	1	0100

Address	Data
000	0101
001	1111
010	0000
011	0110
100	1000
101	0001
110	1010
111	0100

COEN 421/6341: Embedded Systems Design

80

80



## Example: Set-Associative Cache

Set-Associative Cache with two sets

Data requests: 001, 010, 011, 100, 101, and 111

Block	Bank 0 tag	Bank 0 data	Bank 1 tag	Bank 1 data
0	01	0000	10	1000
1	10	0001	11	0100

Address	Data
000	0101
001	1111
010	0000
011	0110
100	1000
101	0001
110	1010
111	0100